

Kaushik Janmanchi

Software Engineer | Cloud Platform & Distributed Systems

 kaushik.janmanchi@gmail.com  8108616917  Hyderabad  linkedin.com/in/kaushik-janmanchi-2978b3157

Results-driven Software & Reliability Engineer with 5.5+ years of experience designing, operating, and automating cloud infrastructure, enterprise data pipelines, and full-stack observability platforms across multi-cloud environments (AWS, GCP). Expert in AWS-native monitoring (CloudWatch, X-Ray), Prometheus, Grafana, Dynatrace, and Terraform-driven infrastructure as code (IaC). Proven track record in 24x7 P1-P4 production incident response, building AI/LLM-assisted diagnostic triage systems, and engineering resilient Python/Java/SQL data ingestion and ETL workflows.

Skills

- **Cloud & Infrastructure:** AWS (CloudWatch, X-Ray, EKS, EC2, S3, RDS, VPC, OpenSearch), GCP (GKE, Compute Engine), Microsoft Azure, Terraform, Terraform Cloud, Helm, Multi-Account Architecture.
- **Observability & SRE:** Prometheus, Grafana, PromQL, Alertmanager, Dynatrace, AWS CloudWatch, Datadog, Splunk, OpenSearch/ELK, SLI/SLO Definitions, Golden Signals, Blameless RCAs, Runbook Automation, ServiceNow ITSM.
- **Data Engineering & Backend:** Python (Advanced Automation, ETL), Java (Spring Boot), SQL, PL/SQL (EXPLAIN PLAN, Query Optimization), Apache Airflow (DAGs), REST APIs, Schema Validation, Data Quality Controls.
- **Containers & DevOps:** Kubernetes (EKS, GKE, AKS), Docker, Podman, ArgoCD (GitOps), GitHub Actions, GitLab CI, Jenkins, Automated Health Checks & Rollbacks.
- **AI & Developer Tooling:** Agentic AI / LLM Tooling (OpenAI API, Bedrock, Cursor, Claude Code, MCP Integrations), Automated Runbook Generation, AI-Assisted Triage.
- **Systems & Networking:** Linux Internals (TCP/IP, Socket Management, JVM Heap/Thread Dumps, I/O Tuning), ACID Databases (Oracle DB, PostgreSQL, MySQL).

Professional Experience

GraphDB Specialist (Technical Support Engineer), Agivant Technologies 10/2025 – Present | Hyderabad

- Serve as the primary technical point of contact for enterprise clients, troubleshooting complex distributed database clusters, network routing, and multi-account cloud infrastructure (AWS/GCP) under strict SLAs.
- Engineered an AI/LLM-powered diagnostic triage assistant that parses incoming logs, correlates CloudWatch/Prometheus telemetry, and auto-suggests runbooks, reducing initial MTTR by 40%.
- Implemented end-to-end cloud observability across AWS and GCP using Prometheus, PromQL, Grafana, and CloudWatch to track p95/p99 latency spikes, pod scheduling bottlenecks, and query execution degradation.
- Developed custom Python and SQL data validation and migration pipelines to execute zero-downtime cross-cloud database transfers with strict schema checks and completeness guarantees.
- Configured automated alerting and threshold-based incident dispatching into ServiceNow/ITSM, authoring comprehensive blameless Root Cause Analyses (RCAs) and runbooks.

Site Reliability Engineer, Oracle Solution Services India Pvt Ltd 06/2024 – 09/2025 | Hyderabad

- Acted as a 24x7 primary on-call responder for mission-critical enterprise SaaS environments serving 22,000+ global customers, maintaining 99.99% system availability.
- Built a parallel Python-based automated diagnostics framework to inspect K8s pod states, JVM heap memory leaks, database connection pool exhaustion, and REST health endpoints with automated self-healing restarts.
- Designed and maintained dynamic multi-stage data consolidation pipelines using Python and Bash, reducing pipeline execution and data sync failures by 65%.
- Managed Infrastructure as Code workflows using Ansible/AWX and Terraform, eliminating configuration drift across global regions (APAC, EMEA, AMER) and cutting operational toil by 45%.
- Diagnosed complex OS and network bottlenecks including Linux TCP socket exhaustion, thread contention, and disk I/O latency.

Associate Software Engineer, Solera Holdings 10/2022 – 06/2024 | Bengaluru

- Engineered high-throughput backend services and ETL data ingestion pipelines in Python and Java, processing large-scale datasets across Oracle DB and MySQL.
- Optimized SQL query execution speeds by 25% through EXPLAIN PLAN analysis, index redesign, and PL/SQL refactoring.
- Built and maintained GitHub Actions and Jenkins CI/CD release pipelines with automated unit tests, schema validation quality gates, and automated deployment rollbacks.
- Implemented structured JSON logging, OpenTelemetry tracing headers, and custom application metrics to expedite defect triage in production environments.

System Trainee & Intern, Intellect Design Arena 02/2021 – 10/2022 | Chennai

- Developed transactional REST microservices using Java, Spring Boot, and Oracle SQL, ensuring strict ACID compliance, data security, and least-privilege access for financial banking gateways.
- Implemented token-based session security (JWT) and encrypted data-in-transit pipelines compliant with financial governance standards.

Education

Master of Computer Applications (MCA), Vivekanand Education Society's Institute of technology 08/2018 – 12/2021 | Mumbai, Maharashtra

Bachelor of Science (Computer Science), J.K. College of Science and Commerce 08/2014 – 11/2017 | Navi Mumbai, Maharashtra

MediFlow AI – Enterprise Healthcare Platform

Tech Stack: AWS (EKS, CloudWatch), GCP (GKE), Kubernetes, Terraform, Helm, ArgoCD, Apache Airflow, Prometheus, Grafana, Python.

- Provisioned multi-region Kubernetes clusters using Terraform and Helm, managing GitOps deployment pipelines via ArgoCD with automated canary rollbacks.
- Built scheduled Apache Airflow ETL DAG workflows to orchestrate batch processing, schema validation, and cross-cluster data synchronization.
- Integrated custom Python webhook listeners with Prometheus Alertmanager and CloudWatch to trigger automated horizontal pod autoscaling and self-healing restarts.

AI Job Radar — AI-Powered Job Discovery & Matching Platform

Tech Stack: Python, FastAPI, Next.js, React, SQLite, SQLAlchemy, APScheduler, OpenAI API, Podman, GitHub Actions, Trivy, GHCR, GCP (GKE, Artifact Registry), Kubernetes, Terraform, Helm.

- Engineered high-concurrency data ingestion pipelines utilizing 24-worker parallel processing to extract, transform, and ingest structured records from 145+ distributed sources.
- Implemented data integrity & deduplication algorithms using cryptographic fingerprinting and schema validation to process irregular payloads with zero data duplication.
- Built a multi-dimensional scoring and evaluation engine implementing async processing, configurable background polling, and persistent state management via SQLAlchemy.
- Containerized microservices using Podman Compose with isolated networking, persistent volumes, and resilient error-handling with deterministic fallbacks.